

Fault Tolerant Resource Adaptor API

Table of Contents

The Fault Tolerant Resource Adaptor Object	1
The Fault Tolerant Resource Adaptor Context	2
The Fault Tolerant Resource Adaptor Replicated Data Sources	2
The Fault Tolerant Resource Adaptor Timer	2

JAIN SLEE Resource Adaptors exist on the boundary between the container and the underlying protocol. The specification contract requires the object to implement the `javax.slee.resource.ResourceAdaptor` interface. This interface defines callbacks, which SLEE uses to interact with the , including one to provide the `javax.slee.resource.ResourceAdaptorContext` object. The Resource Adaptor Context provides object facilities to interact with SLEE.

The JAIN SLEE 1.1 RA API is a major milestone, standardizing RA and JSLEE contract. However, it misses an API for clustering, which is critical for a RA deployed in a clustered JAIN SLEE environment. The JAIN SLEE 1.1 contract does not define any fault tolerant data source nor cluster state callbacks.

The **Fault Tolerant RA API** extends the JAIN SLEE 1.1 RA , providing missing features related to clustering. An effort has been made keep the API similar to the standard RA contract, so that anyone who has developed a JAIN SLEE 1.1 RA is able to easily use the proprietary API extension.

The Fault Tolerant Resource Adaptor Object

The core of the Fault Tolerant RA API is the `org.mobicens.slee.resource.cluster.FaultTolerantResourceAdaptor` interface. It is intended to be used instead of the `javax.slee.resource.ResourceAdaptor` interface from the JAIN SLEE 1.1 Specification.

The FaultTolerant interface provides three new callback methods used by the container:

`setFaultTolerantResourceAdaptorContext(FaultTolerantResourceAdaptorContext context)`

This method provides the RA with the `org.mobicens.slee.resource.cluster.FaultTolerantResourceAdaptorContext` object, which gives access to facilities related with the cluster. This method is invoked by SLEE after invoking `raConfigure(ConfigProperties properties)` from JAIN SLEE 1.1 specs.

`unsetFaultTolerantResourceAdaptorContext()`

This method indicates that the RA should remove any references it has to the `FaultTolerantResourceAdaptorContext`, as it is not valid anymore. The method is invoked by SLEE before invoking `unsetResourceAdaptorContext()` from JAIN SLEE 1.1 specs.

`failOver(K key)`

Callback from SLEE when the local RA was selected to recover the state for a replicated data key, which was owned by a cluster member that failed. The RA may then restore any runtime resources associated with such data.

`dataRemoved(K key)`

Optional callback from SLEE when the replicated data key was removed from the cluster, this may be helpful when the local RA maintains local state.

The Fault Tolerant Resource Adaptor Context

The clustered RA context follows the contract of JAIN SLEE 1.1 specification interface `javax.slee.resource.ResourceAdaptorContext`. It gives access to facilities that the RA may use when run in a clustered environment.

The cluster contract is defined in: `org.mobicens.slee.resource.cluster.FaultTolerantResourceAdaptorContext`. It provides critical information, such as if SLEE is running in local mode (not clustered), if it is the head/master member of the cluster, and what the members of the cluster are.

The Fault Tolerant Resource Adaptor Replicated Data Sources

The Fault Tolerant Resource Adaptor Context provides two data sources to replicate data in cluster:

`ReplicatedData`

A container for serializable data, which is replicated in the SLEE cluster, but don't require any failover.

`ReplicatedDataWithFailover`

A `ReplicatedData` which requires fail over callbacks, this means, that for all data stored here, when a cluster member goes down, the SLEE in another cluster member will invoke the `failOver(Key k)` callback in the Fault Tolerant RA object.

When retrieved from the context through a boolean parameter, both types of `ReplicatedData` can activate the callback on the `FaultTolerantResourceAdaptor` which indicates that a specific data was removed from the cluster remotely.

The Fault Tolerant Resource Adaptor Timer

The standard Resource Adaptor Context provides a `java.util.Timer`, which can be used by the Resource Adaptor to schedule the execution of tasks, the Fault Tolerant Resource Adaptor Context provides `org.mobicens.slee.resource.cluster.FaultTolerantTimer`, an alternative scheduler which is able to fail over tasks scheduled.

The Fault Tolerant Timer has an interface that resembles the JDK's `ScheduledExecutorService`, with two fundamental changes to allow a proper interaction in a cluster environment:

Task Interface

Instead of relying on pure `Runnable` tasks, tasks must follow a specific interface `FaultTolerantTimerTask`, to ensure that the timer is able to replicate the task's data, and failover the task in any cluster node.

Task Cancellation

Cancellation of task is done through the Timer interface, not through `ScheduledFuture` objects,

this allows the operation to be easily done in any cluster node.

The Fault Tolerant Timer interface:

`cancel(Serializable taskID)`

Requests the cancellation of the FT Timer Task with the specified ID.

`configure(FaultTolerantTimerTaskFactory taskFactory, int threads)`

Configures the fault tolerant timer, specifying the timer task factory and the number of threads the timer uses to execute tasks.

`isConfigured()`

Indicates if the timer is configured.

`schedule(FaultTolerantTimerTask task, long delay, TimeUnit unit)`

Creates and executes a one-shot action that becomes enabled after the given delay.

`scheduleAtFixedRate(FaultTolerantTimerTask task, long initialDelay, long period, TimeUnit unit)::` Creates and executes a periodic action that becomes enabled first after the given initial delay, and subsequently with the given period; that is executions will commence after `initialDelay` then `initialDelay+period`, then `initialDelay + 2 * period`, and so on. If any execution of the task encounters an exception, subsequent executions are suppressed. Otherwise, the task will only terminate via cancellation or termination of the executor. If any execution of this task takes longer than its period, then subsequent executions may start late, but will not concurrently execute.

`scheduleWithFixedDelay(FaultTolerantTimerTask task, long initialDelay, long delay, TimeUnit unit)::` Creates and executes a periodic action that becomes enabled first after the given initial delay, and subsequently with the given delay between the termination of one execution and the commencement of the next. If any execution of the task encounters an exception, subsequent executions are suppressed. Otherwise, the task will only terminate via cancellation or termination of the executor.



There is a single Fault Tolerant Timer per RA Entity, and when first retrieved, and before any task can be scheduled, the Fault Tolerant Timer must be configured, through its `configure(...)` method.

The Fault Tolerant Resource Adaptor Timer Task

As mentioned in previous section, tasks submitted to the Fault Tolerant Timer must follow a specific interface, `FaultTolerantTimerTask`, it is nothing more than a `Runnable`, which provides the replicable `FaultTolerantTimerTaskData`. The task data must be serializable and provide a Serializable task ID, which identifies the task, and may be used to cancel its execution.

The Fault Tolerant Resource Adaptor Timer Example Usage

A simple example for the usage of the Fault Tolerant Timer Task:

```
// data, task and factory implementation  
  
package org.mobicenss.slee.resource.sip11;
```

```

import java.io.Serializable;

import org.mobicens.slee.resource.cluster.FaultTolerantTimerTaskData;

public class FaultTolerantTimerTaskDataImpl implements
    FaultTolerantTimerTaskData {

    private final String taskID;

    public FaultTolerantTimerTaskDataImpl(String taskID) {
        this.taskID = taskID;
    }

    @Override
    public Serializable getTaskID() {
        return taskID;
    }
}

package org.mobicens.slee.resource.sip11;

import org.mobicens.slee.resource.cluster.FaultTolerantTimerTask;
import org.mobicens.slee.resource.cluster.FaultTolerantTimerTaskData;
import org.mobicens.slee.resource.cluster.FaultTolerantTimerTaskFactory;

public class FaultTolerantTimerTaskFactoryImpl implements
    FaultTolerantTimerTaskFactory {

    private final SipResourceAdaptor ra;

    public FaultTolerantTimerTaskFactoryImpl(SipResourceAdaptor ra) {
        this.ra = ra;
    }

    @Override
    public FaultTolerantTimerTask getTask(FaultTolerantTimerTaskData data) {
        return new FaultTolerantTimerTaskImpl(ra, data);
    }
}

package org.mobicens.slee.resource.sip11;

import org.mobicens.slee.resource.cluster.FaultTolerantTimerTask;
import org.mobicens.slee.resource.cluster.FaultTolerantTimerTaskData;

public class FaultTolerantTimerTaskImpl implements FaultTolerantTimerTask {

    private final SipResourceAdaptor ra;
    private final FaultTolerantTimerTaskData data;

```

```

    public FaultTolerantTimerTaskImpl(SipResourceAdaptor ra,
        FaultTolerantTimerTaskData data) {
        this.ra = ra;
        this.data = data;
    }

    @Override
    public void run() {
        ra.getTracer("FaultTolerantTimerTaskImpl").info("Timer executed.");
    }

    @Override
    public FaultTolerantTimerTaskData getTaskData() {
        return data;
    }
}

// ra code retrieving the timer, configuring it and submitting a task

public void setFaultTolerantResourceAdaptorContext(
    FaultTolerantResourceAdaptorContext<SipActivityHandle, String> context) {
    this.ftRaContext = context;
    FaultTolerantTimer timer = context.getFaultTolerantTimer();
    timer.config(new FaultTolerantTimerTaskFactoryImpl(this), 4);
    FaultTolerantTimerTaskDataImpl data = new FaultTolerantTimerTaskDataImpl("xyz");
    FaultTolerantTimerTaskImpl task = new FaultTolerantTimerTaskImpl(this,
        data);
    timer.schedule(task, 30, TimeUnit.SECONDS);
}

```